

EgoX: Egocentric Video Generation from a Single Exocentric Video

Anonymous CVPR submission

Paper ID 3508

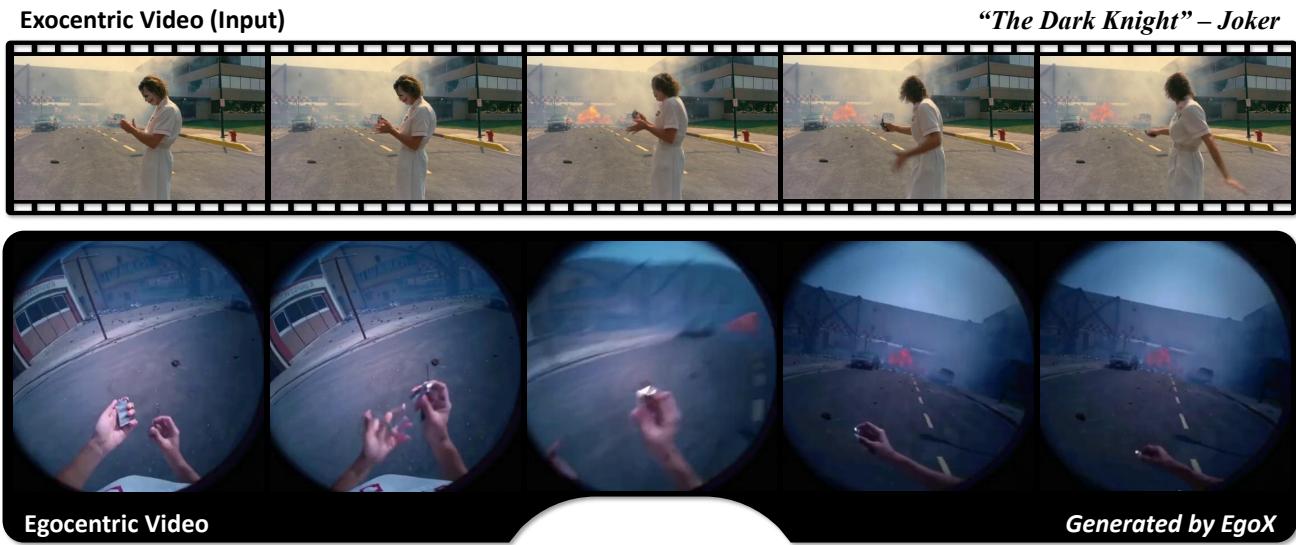


Figure 1. Given a single exocentric video, **EgoX** generates what the scene would look like from the actor’s eyes. Shown with an in-the-wild clip from *The Dark Knight*, our approach achieves realistic and generalizable egocentric generation.

Abstract

001 *Egocentric perception enables humans to experience and*
 002 *understand the world directly from their own point of view.*
 003 *Translating exocentric (third-person) videos into egocentric*
 004 *(first-person) videos opens up new possibilities for immer-*
 005 *sive understanding but remains highly challenging due to*
 006 *extreme camera pose variations and minimal view overlap.*
 007 *This task requires faithfully preserving visible content while*
 008 *synthesizing unseen regions in a geometrically consistent*
 009 *manner. To achieve this, we present **EgoX**, a novel frame-*
 010 *work for generating egocentric videos from a single exocen-*
 011 *tic input. EgoX leverages the pretrained spatio-temporal*
 012 *knowledge of large-scale video diffusion models through*
 013 *lightweight LoRA adaptation and introduces a unified con-*
 014 *ditioning strategy that combines exocentric and egocentric*
 015 *priors via width- and channel-wise concatenation. Addi-*
 016 *tionally, a geometry-guided self-attention mechanism selec-*
 017 *tively attends to spatially relevant regions, ensuring geo-*
 018 *metric coherence and high visual fidelity. Our approach*

achieves coherent and realistic egocentric video generation while demonstrating strong scalability and robustness across unseen and in-the-wild videos.

019
020
021

1. Introduction

022 Don’t you wish you could experience iconic scenes from films like *The Dark Knight* as if you were the *Joker* yourself? Exocentric-to-egocentric video generation makes this possible by converting a third-person scene into a realistic first-person perspective. This capability opens up new possibilities in the film industry, where viewers are no longer limited to passively watching a scene but can step into it and become the main character. They can become a superhero themselves or experience what it is like to play on the field as an MLB player. Beyond entertainment, egocentric perspectives are crucial in fields such as robotics and AR/VR, where understanding how the world appears from the actor’s point of view enables better imitation, reasoning, and interaction [15, 21]. This stems from the fact that humans

023
024
025
026
027
028
029
030
031
032
033
034
035
036

037 perceive and interact with the world through a first-person,
038 egocentric viewpoint.

039 However, generating such first-person perspectives is
040 challenging, since the model must maintain scene consis-
041 tency across views by reconstructing visible areas and re-
042 alistically synthesizing unseen regions. A straightforward
043 way to achieve this is to use a camera control model. Re-
044 cent advances in camera control video generation mod-
045 els [18, 35, 47] have shown impressive performance in gen-
046 erating consistent views under moderate pose variations.
047 However, these methods primarily focus on modest view-
048 point changes, whereas exocentric-to-egocentric video gen-
049 eration requires extreme camera pose translation that dras-
050 tically alters the visible field of view. This difference intro-
051 duces two major challenges. First, extreme viewpoint shifts
052 result in large unseen regions that must be plausibly synthe-
053 sized based on scene understanding rather than direct obser-
054 vation. Second, only a small portion of the exocentric view
055 corresponds to the egocentric perspective, making it crucial
056 for the model to distinguish between view-related informa-
057 tion that should be used as conditioning and unrelated con-
058 tent that should be suppressed. As illustrated in Fig. 2, ef-
059 fective generation therefore requires selectively attending to
060 meaningful regions while discarding irrelevant background
061 areas and plausibly synthesizing uninformed regions in a
062 geometrically consistent manner. Therefore existing cam-
063 era control models do not account for these challenges and
064 thus often fail in exocentric-to-egocentric video generation.

065 Due to the inherent difficulty of this task, previous ap-
066 proaches often avoid generating the egocentric view from
067 scratch or require additional inputs to simplify the problem.
068 EgoExo-Gen [43] takes both an exocentric video and the
069 first egocentric frame as inputs to generate only the subse-
070 quent sequence. Exo2Ego-V [26] utilizes four simultaneous
071 exocentric camera views to capture richer spatial context
072 and reduce the uninformed regions.

073 To address the limitations of previous approaches, we
074 propose EgoX, a novel framework that generates egocen-
075 tric video from a single exocentric video, achieving prac-
076 tical and generalizable egocentric generation from a sin-
077 gle exocentric input. Our method leverages the pretrained
078 spatio-temporal knowledge of large-scale video diffusion
079 models with minimal modification, enabling the model to
080 plausibly synthesize unseen regions in a geometrically con-
081 sistent manner. Specifically, we design a unified condition-
082 ing strategy that combines exocentric views and egocen-
083 tric priors through width-wise and channel-wise integration
084 with clean latent representations, requiring only lightweight
085 LoRA-based adaptation. Furthermore, a geometry-guided
086 self-attention allows the model to focus on spatially relevant
087 regions while suppressing unrelated areas, leading to coher-
088 ent and high-fidelity egocentric video generation. By effec-
089 tively leveraging pretrained weights, our approach produces

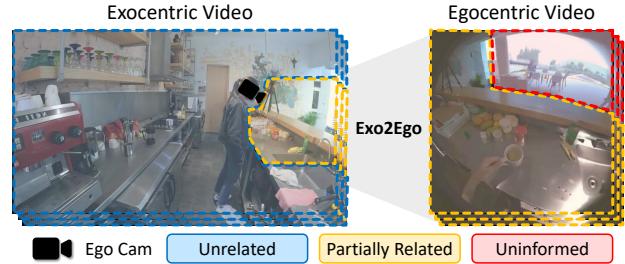


Figure 2. **Exo-to-Ego view generation example.** The model has to preserve view-related content from the exocentric input, generate uninformed regions realistically, and ignore unrelated areas for consistent egocentric synthesis.

high-quality egocentric videos and demonstrates strong
090 generalization across diverse environments, including chal-
091 lenging in-the-wild scenarios, as illustrated in Fig. 1.

To summarize, the major contributions of our paper are
093 as follows:

- We propose a novel framework **EgoX** for synthesizing high-fidelity egocentric video from a *single* exocentric video by effectively exploiting the pretrained video diffusion models.
- We design a unified conditioning strategy that jointly combines exocentric video and egocentric priors through width-wise and channel-wise integration, achieving robust geometric consistency and high-quality generation.
- We introduce a geometry-guided self-attention and clean latent representations that selectively focuses on view-relevant regions and enhances accurate reconstruction, leading to more coherent egocentric synthesis.
- Extensive qualitative and quantitative experiments demonstrate that **EgoX** outperforms previous approaches by a large margin, achieving *state-of-the-art* performance on diverse and challenging exo-to-ego video generation benchmarks.

2. Related Work

2.1. Exo-to-Ego View Generation

Prior works on exo-to-ego view generation have explored various conditioning mechanisms and task formulations to bridge the significant viewpoint gap. Some approaches [26, 27, 31] incorporate exocentric features by concatenating them channel-wise with the egocentric representation. However, this method struggles with the fundamental lack of pixel-wise correspondence between the two viewpoints. This spatial misalignment makes it difficult for the model to effectively leverage the conditioning information, often leading to a poor understanding of the scene geometry, which can result in overfitting or a degradation in output quality. Other works, such as 4Diff [10], employ cross-attention mechanisms to condition the generation on exo-

127 centric views. This approach, however, prevents the utilization
 128 of powerful pretrained diffusion weights, limiting its
 129 generalizability and resulting in lower-quality synthesis.

130 To address these limitations, other methods utilize ref-
 131 erence frames or multi-view conditions. For instance,
 132 EgoExo-Gen [43] require the first egocentric frame to gen-
 133 erate the rest of the sequence. Exo2Ego-V [26] performs
 134 full video translation but relies on four exocentric video in-
 135 puts and separately trained spatial and temporal modules,
 136 which limits its generalization and fails to fully exploit spa-
 137 tio-temporal priors. In contrast, our model generalizes ef-
 138 fectively using pretrained video diffusion weights while re-
 139 quiring only a single exocentric input.

140 2.2. Video Diffusion Models

141 Recent advancements in video diffusion models [1, 5, 6,
 142 14, 39, 45] have led to significant improvements in gener-
 143 ative quality, producing highly realistic and coherent video
 144 sequences. This has spurred a wide range of research ex-
 145 ploring how to utilize these powerful generative capabili-
 146 ties in various applications [9, 19, 20, 32, 49]. A key area
 147 of this research focuses on conditional video generation,
 148 where the synthesis process is guided by specific inputs.
 149 Many works [7, 19, 22, 44, 49] have demonstrated success-
 150 ful control using conditions such as depth maps or static
 151 images.

152 Building on this, several methods have been proposed for
 153 camera-controlled video generation [4, 28, 47]. These ap-
 154 proaches can be broadly categorized into two main groups.
 155 The first group [3, 4, 29, 42, 46] conditions the diffusion
 156 model directly on camera extrinsic parameters, often repre-
 157 sented as raw matrices or Plücker coordinates. The second
 158 group [18, 25, 35, 41, 47] first lifts the input video into an
 159 intermediate 3D representation, such as a point cloud. This
 160 3D scene is then rendered from a new, user-specified cam-
 161 era pose, and the resulting image is used as a strong spatial
 162 condition to guide the final video generation.

163 However, existing methods for camera control are
 164 primarily designed for modest changes in viewpoint.
 165 They struggle to handle the extreme camera pose differ-
 166 ences, a challenge that becomes particularly significant in
 167 exocentric-to-egocentric video generation. Our work ad-
 168 dresses this critical gap by proposing a model capable of
 169 generating coherent egocentric videos from a significantly
 170 different exocentric perspective.

171 3. Method

172 Given an exocentric video sequence $X = \{X_i\}_{i=0}^F$ and ego-
 173 centric camera pose $\phi = \{\phi_i\}_{i=0}^F$, the goal is to generate
 174 a corresponding egocentric video sequence $Y = \{Y_i\}_{i=0}^F$
 175 that depicts the same scene from a first-person viewpoint.
 176 The key challenge is to preserve the visible content in the
 177 exocentric view while synthesizing unseen regions in a ge-

178 ometrically consistent and realistic manner. To this end, the
 179 exocentric sequence X is first lifted into a 3D represen-
 180 tation and rendered from the target egocentric viewpoint
 181 (Sec. 3.1), which becomes an egocentric prior video P .
 182 Both P and the original exocentric video X are then pro-
 183 vided as inputs to a video diffusion model (Sec. 3.2). In
 184 addition, a geometry-guided self-attention (Sec. 3.3) is pro-
 185 posed to adaptively focus on view-consistent regions and
 186 enhance feature coherence across perspectives.

187 3.1. Egocentric Point Cloud Rendering

188 For this stage, we render an egocentric prior video $P \in$
 $189 \mathbb{R}^{F \times 3 \times H \times W}$ via point cloud rendering from the exocen-
 190 tric view. This prior provides both explicit pixel-wise
 191 RGB information and implicit camera trajectory cues that
 192 guide viewpoint alignment. Specifically, we first estimate a
 193 monocular depth map $D^m \in \mathbb{R}^{F \times H \times W}$ for each frame us-
 194 ing a single-image depth estimator [40], and a video-based
 195 depth map $D^v \in \mathbb{R}^{F \times H \times W}$ using a temporal depth estima-
 196 tor [8]. Because D^m is estimated independently per frame,
 197 depth values often exhibit slight inconsistencies across time.
 198 In contrast, D^v produces a temporally smooth yet affine-
 199 invariant depth estimate. To combine the advantages of
 200 both, we temporally align D^v with D^m . Following [16],
 201 we optimize affine transformation parameters α, β using a
 202 momentum-based update strategy, yielding $\hat{\alpha} = \{\hat{\alpha}_f\}_{f=0}^F$
 203 and $\hat{\beta} = \{\hat{\beta}_f\}_{f=0}^F$, which represent the per-frame affine
 204 transformations. The final aligned depth is computed as:

$$D^f = \frac{1}{\hat{\alpha}/D^v + \hat{\beta}}, \quad (1)$$

205 where D^f denotes the final aligned depth map. Dynamic
 206 objects are masked out so that only static background re-
 207 gions are used during both alignment and rendering. For
 208 further details, please refer to [16].

209 After obtaining the aligned depth map D^f , we convert it
 210 into a 3D point cloud representation using the correspond-
 211 ing camera intrinsics. We then render the egocentric prior
 212 frames using a point cloud renderer [33]:

$$P = \text{render}(X, D^f, \phi), \quad (2)$$

213 where $X \in \mathbb{R}^{F \times 3 \times H \times W}$ is the exocentric RGB video and
 ϕ is egocentric camera poses.

214 3.2. Exo-to-Ego View Generation with VDM

215 As illustrated in Fig. 3, the model takes an exocentric
 216 video $X \in \mathbb{R}^{F \times 3 \times H \times W}$ and the egocentric prior video
 217 $P \in \mathbb{R}^{F \times 3 \times H \times W'}$ as conditioning inputs. Both inputs are
 218 encoded by a frozen VAE encoder, producing latent fea-
 219 tures $x_0 \in \mathbb{R}^{f \times c \times h \times w}$ and $p_0 \in \mathbb{R}^{f \times c \times h \times w'}$, respectively.
 220 These latents are then concatenated with the noisy latent
 221 $z_t \in \mathbb{R}^{f \times c \times h \times w'}$ to form the input of the diffusion model.

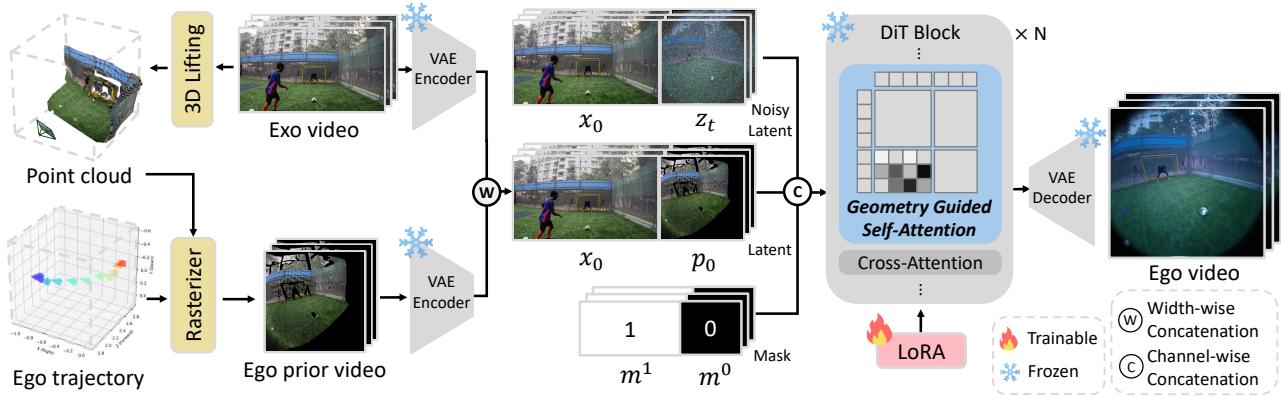


Figure 3. **Overall pipeline.** Given an exocentric video input, we first lift it into a 3D point cloud and render the scene from the egocentric viewpoint to obtain the egocentric prior video. The clean exocentric video latent and the egocentric prior latent are combined via width-wise and channel-wise concatenation in the latent space, and then fed into a pretrained video diffusion model equipped with the proposed geometry-guided self-attention.

The egocentric prior latent p_0 shares the same viewpoint as the target egocentric video and therefore preserves pixel-wise correspondence. We concatenate p_0 with z_t along the channel dimension, providing viewpoint-aligned and temporally coherent guidance during generation. Although p_0 offers explicit geometric cues for the regions visible in the rendered ego view, it remains noisy and lacks substantial portions of the scene. To complement the missing information in the rendered egocentric view, we further use the exocentric video latent x_0 to provide broader scene context. Since the viewpoint of x_0 differs from that of the noisy egocentric latent z_t , their features are not pixel-wise aligned. Therefore, we concatenate x_0 with z_t along the width dimension, encouraging the model to infer cross-view correspondences and perform spatial warping implicitly. Unlike [17], which utilizes SDEdit [30] by concatenating a noisy conditioning latent with a noisy target latent for conditional generation, our method concatenates the clean latent x_0 with the noisy z_t throughout all denoising timesteps, while only z_t is updated and x_0 remains fixed. This design encourages the model to consistently reference fine-grained details from x_0 , enabling more accurate and reliable spatial warping.

The overall relation between inputs and outputs is defined as:

$$z_{t-1} = f_\theta(x_0, z_t | x_0, p_0 | m^1, m^0), \quad (3)$$

where f_θ denotes a single-step denoising function of the VDM, x_0 is the exocentric video latent, p_0 is the egocentric prior latent, and m is the binary mask specifying whether each spatial region is used for conditioning or for synthesis. Once the sampling is complete, we remove the exocentric part of the latent and decode only the egocentric part to obtain the final result.

3.3. Geometry-Guided Self-Attention

As mentioned in Sec. 1, the exocentric video condition includes irrelevant regions that can distract the model during exo-to-ego view generation. To address this, we introduce a Geometry-Guided Self-Attention (GGA) that adaptively emphasizes spatially corresponding regions between exocentric and egocentric representations. When egocentric query tokens $q_{\text{ego}} \in \mathbb{R}^{l \times c}$ attend to exocentric key tokens $k_{\text{exo}} \in \mathbb{R}^{l' \times c}$, the attention should jointly account for semantic similarity (i.e., appearance) and 3D spatial alignment. Ideally, tokens that are both semantically similar and geometrically aligned with the egocentric viewpoint should receive higher attention weights, while unrelated or misaligned regions are suppressed to ensure geometric consistency and realism in the generated views.

To achieve this, we leverage self-attention augmentation with 3D geometric cues. Using the 3D point cloud obtained in Sec. 3.1, we compute 3D direction vectors from the ego camera centers $c = \{c_i\}_{i=0}^F$, $c_i \in \mathbb{R}^3$ in world space to each query and key token position, $\tilde{q}, \tilde{k} \in \mathbb{R}^3$. The unit direction vectors are defined as $\hat{q} = \frac{\tilde{q} - c_i}{\|\tilde{q} - c_i\|_2}$, $\hat{k} = \frac{\tilde{k} - c_i}{\|\tilde{k} - c_i\|_2}$. We then compute the cosine similarity between the two direction vectors and incorporate it into the attention computation as a multiplicative geometric prior.

Specifically, the modified attention logits are formulated as:

$$s'_{m,n} = s_{m,n} + \log(g(\hat{q}_m, \hat{k}_n) \cdot \lambda_g), \quad (4)$$

$$g(\hat{a}, \hat{b}) = \text{cos_sim}(\hat{a}, \hat{b}) + 1, \quad (5)$$

where $s_{m,n} = \frac{q_m^\top k_n}{\sqrt{c}}$ denotes the standard attention logits [38] and λ_g is a hyperparameter that balances this geometry bias term defined in Eq. (5). We add one to the cosine similarity term to ensure positive values before taking the logarithm.

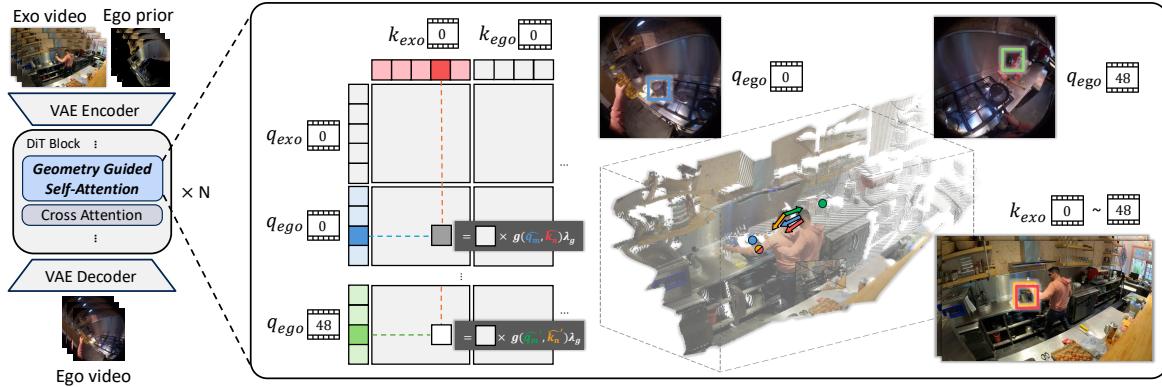


Figure 4. Geometry-Guided Self-Attention Overview. 3D direction similarities between egocentric queries and exocentric keys are used as an additive bias in the attention map, guiding the model to focus on geometrically aligned regions. Although the orange and red directions are the same key tokens, their directions differ due to different camera centers. The blue-red pairs have similar directions and thus receive higher scores, whereas the green-orange pairs have opposite directions and obtain lower scores.

Finally, given an egocentric query q_m and an exocentric key k_n , the attention weight $a_{m,n}$ is computed as:

$$a_{m,n} = \frac{\exp(s'_{m,n})}{\sum_{j=1}^l \exp(s'_{m,j})} \quad (6)$$

$$= \frac{\exp(s_{m,n}) g(\hat{q}_m, \hat{k}_n) \lambda_g}{\sum_{j=1}^l \exp(s_{m,j}) g(\hat{q}_m, \hat{k}_j) \lambda_g}. \quad (7)$$

This formulation allows the attention mechanism to be explicitly guided by geometric alignment between query and key directions, improving spatial consistency and visual coherence across views.

In image generation, spatial relationships can be encoded by multiplying rotation matrices to each query and key before attention, as done in [10, 23, 24, 37]. However, in video generation, the camera center of q_{ego} changes at every frame, making it necessary to compute key directions relative to each query separately. This implies that the geometry bias term should be recomputed for every query–key pair within each frame’s attention operation. As illustrated in Fig. 4, even k_{exo} located at the same position (e.g. red) may have entirely different direction vectors (e.g. red and orange) depending on the camera pose. To handle this, we compute all pairwise direction similarities between k_{exo} and q_{ego} and use this term as an additive bias attention mask, allowing us to reuse optimized attention kernels. This formulation provides a precise geometry-guided self-attention that effectively aligns exocentric and egocentric representations.

4. Experiments

In the following sections, we aim to answer the following research questions that guide our experimental evaluation:

- How does our method outperform existing baselines in both qualitative and quantitative evaluations? (Sec. 4.2,

Sec. 4.3)

- How accurately does the model reconstruct regions visible in the exocentric view? (Sec. 4.1, Sec. 4.3)
- How well does the model generalize to unseen scenes and challenging in-the-wild videos? (Sec. 4.2, Sec. 4.3)
- How does each proposed component contribute to overall performance and generation quality? (Sec. 4.4)

4.1. Experimental Setup

Implementation Details. To support channel-wise concatenation of noisy latent and ego prior latent, we adopt the inpainting variant of Wan 2.1 (14B) Image-to-Video model [39] as our base model. We fine-tuned the model using LoRA (rank = 256) with a batch size of 1, and a single day on 8 H200 (140 GB) GPUs. For the dataset, we curated 4,000 clips from Ego-Exo4D [12] covering diverse scenes and actions, using 3,600 clips for training and 400 for testing. Additionally, we collected 100 unseen clips that are not included in the training set to evaluate generalization performance. More detailed information can be found in Sec. F.

Baselines. Among existing exocentric-to-egocentric video generation approaches, Exo2Ego-V [26] and EgoExo-Gen [43] serve as representative baselines. We adopt Exo2Ego-V as our primary baseline, as EgoExo-Gen does not provide publicly available implementation. With the rapid progress in conditional video generation and camera control models, several recent methods have demonstrated performance comparable to or even surpassing Exo2Ego-V. Therefore, we additionally included Trajectory Crafter [47], a state-of-the-art camera control model, as well as Wan Fun Control [2] and Wan VACE [19], which offer distinct conditioning approach. Wan Fun Control applies channel-wise concatenation for conditioning, and Wan VACE employs an auxiliary conditioning network,

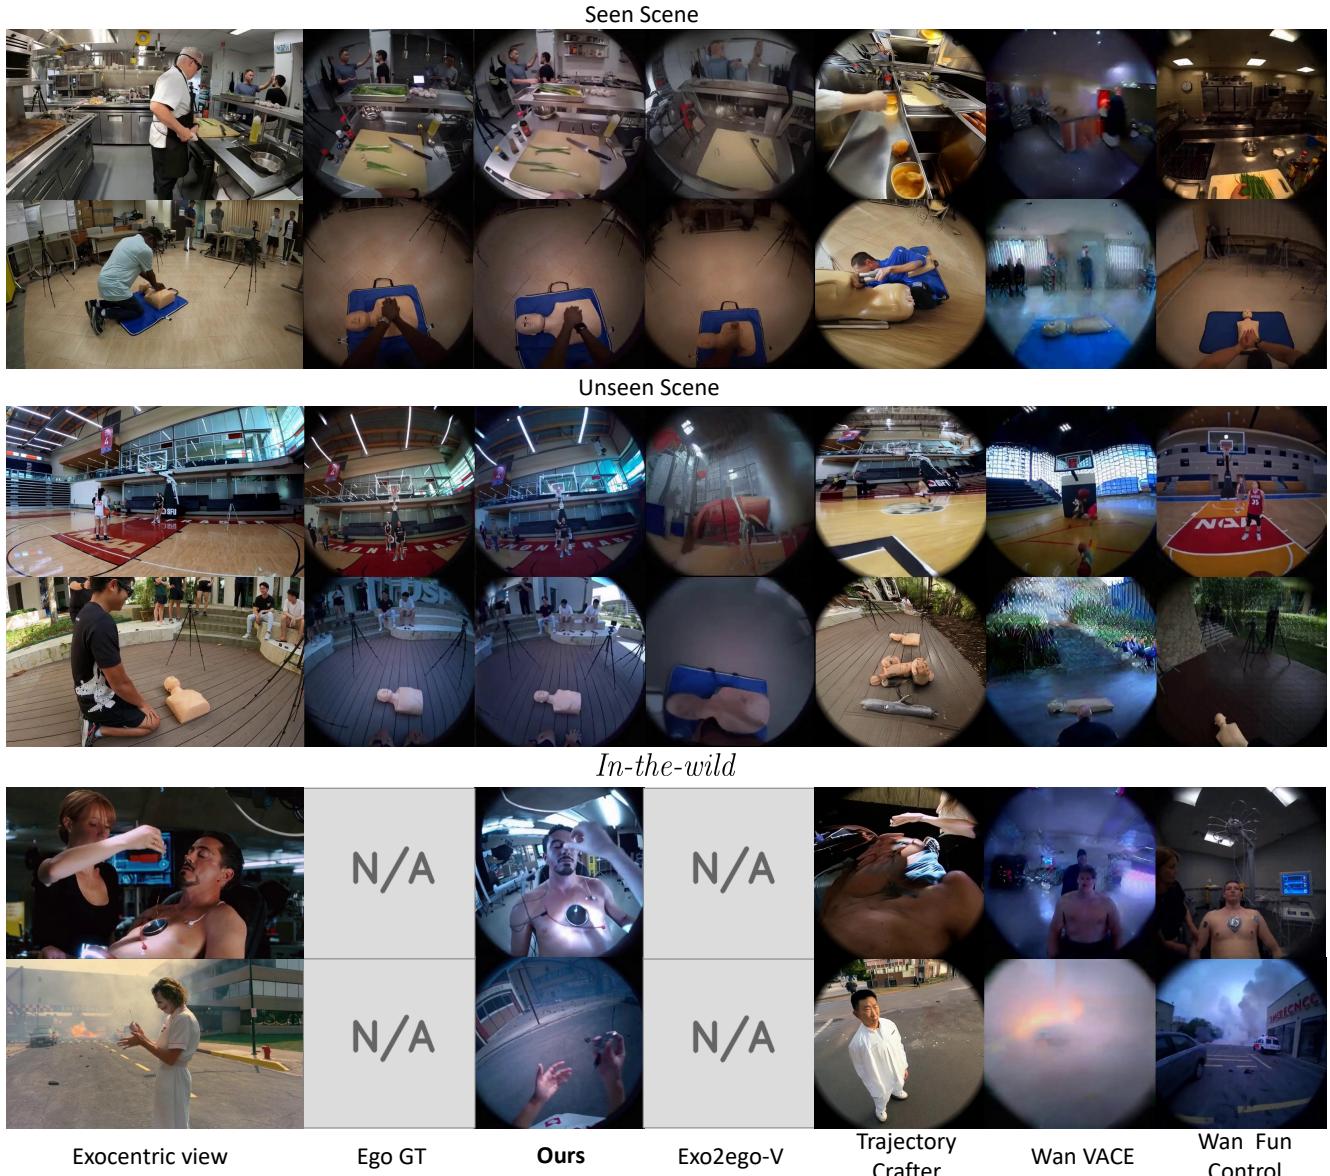


Figure 5. **Qualitative comparison.** Each example shows the exocentric input views and the corresponding generated egocentric views. While other methods fail to reconstruct realistic and coherent videos, our approach produces geometrically accurate and high-quality egocentric generations. N/A indicates that the result is unavailable either due to missing ground truth or the need for additional input views.

354 providing diverse points of comparison for our method. For
355 the fair comparison, we finetuned these baselines using the
356 same training dataset as ours.

357 **Evaluation Metrics.** To evaluate the quality of generated
358 videos, we employed three types of criteria.

- 359 • **Image Criteria.** We measured PSNR, SSIM, LPIPS,
360 and CLIP-I to assess how closely each generated frame
361 matches the ground-truth distribution.
- 362 • **Object Criteria.** Following the object-level evaluation
363 protocol of Ego-Exo4D [13], we assessed object-level

364 consistency between the generated egocentric video and
365 the ground truth. We used SAM2 [34] to segment and
366 track objects and DINOv3 [36] to establish correspondences.
367 For each matched object, we evaluated center-
368 location error, Intersection-Over-Union(IoU), and Con-
369 tour Accuracy to measure spatial alignment and bound-
370 ary fidelity.

- 371 • **Video Criteria.** We measured FVD [11] to evaluate
372 how closely the generated video aligns with the ground-
373 truth distribution. In addition, we assessed VBench [48]-
374 Temporal Flickering, Motion Smoothness, and Dynamic
375 Degree to quantify temporal stability and motion quality.

Scenarios	Method	Image Criteria				Object Criteria			Video Criteria			
		PSNR ↑	SSIM ↑	LIPIS ↓	CLIP-I ↑	Location Error ↓	IoU ↑	Contour Accuracy ↑	FVD ↓	Temporal Flickering ↑	Motion Smoothness ↑	Dynamic Degree ↑
Seen Scenes	Exo2Ego-V	<u>14.53</u>	0.384	<u>0.569</u>	0.774	156.66	0.074	0.364	622.47	0.960	0.966	0.985
	TrajectoryCrafter	13.05	0.375	0.606	0.780	<u>100.74</u>	<u>0.128</u>	<u>0.427</u>	546.09	0.960	0.980	0.947
	Wan Fun Control	12.25	<u>0.463</u>	0.617	0.810	112.57	0.076	0.417	595.07	0.968	0.980	0.901
	Wan VACE	12.95	0.413	0.626	<u>0.829</u>	109.62	0.114	0.376	508.69	0.989	0.994	0.673
	EgoX (Ours)	16.05	0.556	0.498	0.896	61.81	0.363	0.546	184.47	0.977	0.990	0.974
Unseen Scenes	Exo2Ego-V	12.70	<u>0.439</u>	<u>0.597</u>	0.679	214.32	0.003	0.296	1283.50	0.971	0.976	<u>0.978</u>
	TrajectoryCrafter	12.24	0.297	0.619	0.778	192.16	0.039	0.301	<u>821.71</u>	0.966	0.984	0.944
	Wan Fun Control	<u>13.59</u>	0.439	0.604	0.799	191.40	0.042	0.329	968.78	0.971	0.985	0.944
	Wan VACE	12.17	0.345	0.638	<u>0.820</u>	191.97	0.038	0.314	1045.45	0.995	0.996	0.427
	EgoX (Ours)	14.38	0.457	0.552	0.877	149.93	0.092	0.481	440.64	0.981	0.992	0.989

Table 1. **Quantitative Results.** Comparison on image, object, and video metrics. Our method achieves the best overall performance, with Wan VACE showing higher video scores due to static outputs. **Best** results are highlighted in bold, and second-best results are underlined.

Method	Image Criteria				Object Criteria			Video Criteria			
	PSNR ↑	SSIM ↑	LIPIS ↓	CLIP-I ↑	Location Error ↓	IoU ↑	Contour Accuracy ↑	FVD ↓	Temporal Flickering ↑	Motion Smoothness ↑	Dynamic Degree ↑
EgoX (Ours)	16.05	0.556	0.498	0.896	61.81	0.363	0.546	184.47	0.977	0.989	0.974
w/o GGA	14.77	<u>0.539</u>	<u>0.530</u>	0.897	<u>64.30</u>	<u>0.326</u>	<u>0.538</u>	254.08	0.969	0.987	<u>0.877</u>
w/o Ego prior	13.67	0.479	0.573	0.864	90.70	0.417	0.464	211.50	<u>0.974</u>	0.990	0.802
w/o clean latent	<u>15.07</u>	0.528	0.540	0.861	70.17	0.376	0.506	343.33	0.963	0.986	0.864

Table 2. **Ablation Study Results.** Performance comparison by removing each core component of our framework. The full model achieves the best results, while excluding geometry-guided self-attention, ego prior, or clean latent conditioning causes performance degradation. **Best** results are highlighted in bold, and second-best results are underlined.

376

4.2. Qualitative Results

377 Fig. 5 visualizes the qualitative comparisons between our
378 method and the baselines. Note that in the *in-the-wild* scenario,
379 ground-truth egocentric videos are unavailable, and Exo2Ego-V is also not applicable since only a single ex-
380 ocentric video is provided, which does not meet its four-
381 view input requirement. Exo2Ego-V fails to generate high-
382 fidelity frames even when using four exocentric inputs,
383 whereas our model achieves superior visual quality and gen-
384 eralizes well to unseen scenes from only a single exocentric
385 view. Trajectory Crafter struggles with large camera trans-
386 lations, producing spatial distortions and temporal incon-
387 sistencies. Both Wan VACE and Wan Fun Control fail to ef-
388 fectively utilize the exocentric conditioning input, resulting
389 in mismatched geometry, degraded realism, and the inclu-
390 sion of irrelevant exocentric content in the egocentric view.
391 Overall, these results demonstrate that our model effectively
392 leverages pretrained video diffusion knowledge to generate
393 geometrically accurate, visually coherent, and highly realis-
394 tic egocentric videos, maintaining strong performance even
395 under challenging *in-the-wild* conditions. More qualitative
396 results, including temporally aligned visualizations, can be
397 found in Sec. H.

398

4.3. Quantitative Results

400 As shown in Tab. 1, our method achieves the best overall
401 performance across both image and object criteria. In partic-
402 ular, we observe a significant performance gap in the object-

based criteria, indicating that our approach preserves scene
403 geometry and object consistency more effectively than other
404 baselines. While image-level scores may appear slightly
405 lower due to the inherent challenge of synthesizing unsee-
406 n regions that differ from the ground-truth egocentric view,
407 our method still achieves the best results across all image
408 metrics. For video-based metrics, Wan VACE records the
409 highest temporal smoothness and flicker scores. However,
410 this is largely attributed to its generation of overly static
411 videos with limited motion, resulting in low dynamic re-
412 alism. In contrast, our model produces temporally coherent
413 and visually dynamic sequences, demonstrating a better bal-
414 ance between spatial fidelity and motion realism.

415

4.4. Ablation Study

416 We conducted ablation studies to evaluate the contribu-
417 tion of each core component in our framework, including
418 the geometry-guided self-attention (GGA), the egocentric
419 prior conditioning, and the clean latent representation. For
420 each ablation variant, one component was removed while
421 keeping all other settings identical. Quantitative evalua-
422 tions were performed on the seen scene subset to ensure a con-
423 trolled comparison. As shown in Fig. 6 and Tab. 2, remov-
424 ing any of these components results in a noticeable perfor-
425 mance drop, both qualitatively and quantitatively. Without
426 GGA, the model fails to maintain geometric alignment, at-
427 tending to broad and unrelated regions, which leads to spa-
428 tial inconsistency. Without the egocentric prior, the model

417

418

419

420

421

422

423

424

425

426

427

428

429

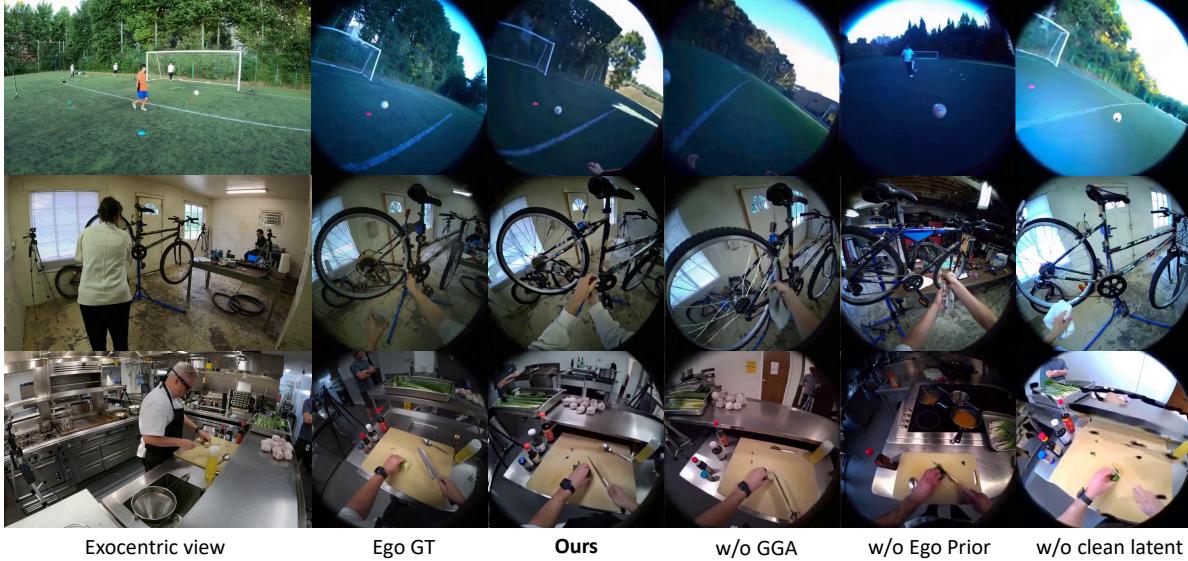


Figure 6. **Ablation qualitative comparison.** Visual results when removing each core component. Removing any single component—GGA, the egocentric prior, or the clean latent representation—results in degraded generation quality and geometric consistency.

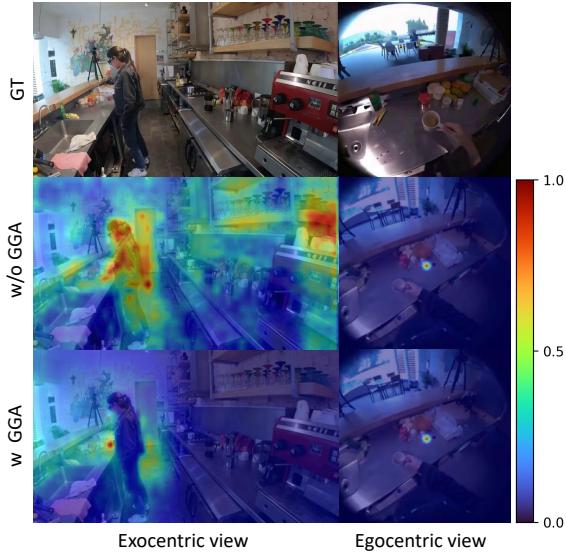


Figure 7. **Attention map visualization.** Visualization of the attention weights when querying the center token of the egocentric view. Without GGA, the model attends to unrelated regions, whereas with GGA, attention is concentrated on related regions, highlighting improved spatial alignment.

lacks explicit pixel-wise and camera trajectory information, thus struggling to follow the correct viewpoint and producing visually implausible frames. Without the clean latent, the exocentric latent is concatenated in a noisy state, which blurs fine-grained details. As a result, the target latent fails to preserve these details, leading to missing or degraded object structures. In the last row of Fig. 6, for instance, the model does not generate the spoon or the small circular in-

gredients on the cutting board that appear in the ground-truth egocentric view.

To further demonstrate the effectiveness of the geometry-guided self-attention, we visualize the attention maps queried by egocentric tokens. As shown in Fig. 7, without GGA, the model attends to broad irrelevant regions, while with GGA, it sharply focuses on view-relevant areas, reinforcing geometric coherence and stabilizing feature alignment. Additional ablation studies are provided in Sec. G.

5. Conclusion

We introduce **EgoX**, the first framework capable of generating egocentric videos from a single exocentric input while achieving strong generalization across diverse scenes. Our method introduces a unified conditioning strategy that combines exocentric and egocentric priors via width- and channel-wise concatenation for effective global context and viewpoint alignment, while leveraging lightweight LoRA-based adaptation to preserve the pretrained video diffusion model’s spatio-temporal reasoning ability. Furthermore, clean latent representations and geometry-guided self-attention enable the model to selectively focus on spatially relevant regions and maintain geometric consistency, resulting in coherent and high-fidelity egocentric generation. Despite its effectiveness, our current framework requires an egocentric camera pose as input. Although this information can be provided interactively by users, incorporating an automatic head-pose estimation module would be a valuable future direction.

438
439
440
441
442
443
444
445
446

447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465

466

References

- [1] Niket Agarwal, Arslan Ali, Maciej Bala, Yogesh Balaji, Erik Barker, Tiffany Cai, Prithvijit Chattopadhyay, Yongxin Chen, Yin Cui, Yifan Ding, et al. Cosmos world foundation model platform for physical ai. *arXiv preprint arXiv:2501.03575*, 2025. 3
- [2] aigc-apps. VideoX-Fun: A flexible framework for video generation. <https://github.com/aigc-apps/VideoX-Fun>, 2024. Accessed: YYYY-MM-DD. 5
- [3] Sherwin Bahmani, Ivan Skorokhodov, Guocheng Qian, Aliaksandr Siarohin, Willi Menapace, Andrea Tagliasacchi, David B. Lindell, and Sergey Tulyakov. Ac3d: Analyzing and improving 3d camera control in video diffusion transformers. *Proc. CVPR*, 2025. 3
- [4] Jianhong Bai, Menghan Xia, Xiao Fu, Xintao Wang, Lianrui Mu, Jinwen Cao, Zuo Zhu Liu, Haoji Hu, Xiang Bai, Pengfei Wan, et al. Recammaster: Camera-controlled generative rendering from a single video. *arXiv preprint arXiv:2503.11647*, 2025. 3
- [5] Andreas Blattmann, Tim Dockhorn, Sumith Kulal, Daniel Mendelevitch, Maciej Kilian, Dominik Lorenz, Yam Levi, Zion English, Vikram Voleti, Adam Letts, et al. Stable video diffusion: Scaling latent video diffusion models to large datasets. *arXiv preprint arXiv:2311.15127*, 2023. 3
- [6] Junsong Chen, Yuyang Zhao, Jincheng Yu, Ruihang Chu, Junyu Chen, Shuai Yang, Xianbang Wang, Yicheng Pan, Daquan Zhou, Huan Ling, et al. Sana-video: Efficient video generation with block linear diffusion transformer. *arXiv preprint arXiv:2509.24695*, 2025. 3
- [7] Liyang Chen, Tianxiang Ma, Jiawei Liu, Bingchuan Li, Zhuowei Chen, Lijie Liu, Xu He, Gen Li, Qian He, and Zhiyong Wu. Humo: Human-centric video generation via collaborative multi-modal conditioning, 2025. 3
- [8] Sili Chen, Hengkai Guo, Shengnan Zhu, Feihu Zhang, Zilong Huang, Jiashi Feng, and Bingyi Kang. Video depth anything: Consistent depth estimation for super-long videos. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 22831–22840, 2025. 3
- [9] Zhaoxi Chen, Tianqi Liu, Long Zhuo, Jiawei Ren, Zeng Tao, He Zhu, Fangzhou Hong, Liang Pan, and Ziwei Liu. 4dnex: Feed-forward 4d generative modeling made easy. *arXiv preprint arXiv:2508.13154*, 2025. 3
- [10] Feng Cheng, Mi Luo, Huiyu Wang, Alex Dimakis, Lorenzo Torresani, Gedas Bertasius, and Kristen Grauman. 4diff: 3d-aware diffusion model for third-to-first viewpoint translation. In *European Conference on Computer Vision*, pages 409–427. Springer, 2024. 2, 5
- [11] Songwei Ge, Aniruddha Mahapatra, Gaurav Parmar, Jun-Yan Zhu, and Jia-Bin Huang. On the content bias in fréchet video distance. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 6
- [12] Kristen Grauman, Andrew Westbury, Eugene Byrne, Zachary Chavis, Antonino Furnari, Rohit Girdhar, Jackson Hamburger, Hao Jiang, Miao Liu, Xingyu Liu, et al. Ego4d: Around the world in 3,000 hours of egocentric video. In

Proceedings of the IEEE/CVF conference on computer vision and pattern recognition, pages 18995–19012, 2022. 5

- [13] Kristen Grauman, Andrew Westbury, Lorenzo Torresani, Kris Kitani, Jitendra Malik, Triantafyllos Afouras, Kumar Ashutosh, Vijay Baiyya, Siddhant Bansal, Bikram Boote, et al. Ego-exo4d: Understanding skilled human activity from first-and third-person perspectives. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 19383–19400, 2024. 6
- [14] Yoav HaCohen, Nisan Chiprut, Benny Brazowski, Daniel Shalem, Dudu Moshe, Eitan Richardson, Eran Levin, Guy Shiran, Nir Zabari, Ori Gordon, et al. Ltx-video: Realtime video latent diffusion. *arXiv preprint arXiv:2501.00103*, 2024. 3
- [15] Yuhang Hu, Boyuan Chen, and Hod Lipson. Egocentric visual self-modeling for autonomous robot dynamics prediction and adaptation. *npj Robotics*, 3(1):14, 2025. 1
- [16] Jiahui Huang, Qunjie Zhou, Hesam Rabeti, Aleksandr Kordova, Huan Ling, Xuanchi Ren, Tianchang Shen, Jun Gao, Dmitry Slepichev, Chen-Hsuan Lin, et al. Vipe: Video pose engine for 3d geometric perception. *arXiv preprint arXiv:2508.10934*, 2025. 3
- [17] Lianghua Huang, Wei Wang, Zhi-Fan Wu, Yupeng Shi, Huanzhang Dou, Chen Liang, Yutong Feng, Yu Liu, and Jingen Zhou. In-context lora for diffusion transformers. *arXiv preprint arxiv:2410.23775*, 2024. 4
- [18] Tianyu Huang, Wangguandong Zheng, Tengfei Wang, Yuhao Liu, Zhenwei Wang, Junta Wu, Jie Jiang, Hui Li, Rynson WH Lau, Wangmeng Zuo, et al. Voyager: Long-range and world-consistent video diffusion for explorable 3d scene generation. *arXiv preprint arXiv:2506.04225*, 2025. 2, 3
- [19] Zeyinzi Jiang, Zhen Han, Chaojie Mao, Jingfeng Zhang, Yulin Pan, and Yu Liu. Vace: All-in-one video creation and editing. *arXiv preprint arXiv:2503.07598*, 2025. 3, 5
- [20] Zeren Jiang, Chuanxia Zheng, Iro Laina, Diane Larlus, and Andrea Vedaldi. Geo4d: Leveraging video generators for geometric 4d scene reconstruction, 2025. 3
- [21] Daekyun Kim, Brian Byunghyun Kang, Kyu Bum Kim, Hyungmin Choi, Jeesoo Ha, Kyu-Jin Cho, and Sungho Jo. Eyes are faster than hands: A soft wearable robot learns user intention from the egocentric view. *Science Robotics*, 4(26):eaav2949, 2019. 1
- [22] Kinam Kim, Junha Hyung, and Jaegul Choo. Temporal in-context fine-tuning for versatile control of video diffusion models. *arXiv preprint arXiv:2506.00996*, 2025. 3
- [23] Xin Kong, Shikun Liu, Xiaoyang Lyu, Marwan Taher, Xiaojuan Qi, and Andrew J Davison. Eschernet: A generative model for scalable view synthesis. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 9503–9513, 2024. 5
- [24] Ruilong Li, Brent Yi, Junchen Liu, Hang Gao, Yi Ma, and Angjoo Kanazawa. Cameras as relative positional encoding. *arXiv preprint arXiv:2507.10496*, 2025. 5
- [25] Teng Li, Guangcong Zheng, Rui Jiang, Shuigen Zhan, Tao Wu, Yehao Lu, Yining Lin, and Xi Li. Realcam-i2v: Real-world image-to-video generation with interactive complex camera control. *arXiv preprint arXiv:2502.10059*, 2025. 3

- 579 [26] Jia-Wei Liu, Weijia Mao, Zhongcong Xu, Jussi Keppo, and
580 Mike Zheng Shou. Exocentric-to-egocentric video genera-
581 tion. *Advances in Neural Information Processing Systems*,
582 37:136149–136172, 2024. 2, 3, 5
- 583 [27] Mi Luo, Zihui Xue, Alex Dimakis, and Kristen Grauman. Put
584 myself in your shoes: Lifting the egocentric perspective from
585 exocentric videos. In *European Conference on Computer
586 Vision*, pages 407–425. Springer, 2024. 2
- 587 [28] Yawen Luo, Jianhong Bai, Xiaoyu Shi, Menghan Xia, Xintao
588 Wang, Pengfei Wan, Di Zhang, Kun Gai, and Tianfan Xue.
589 Camclonemaster: Enabling reference-based camera control
590 for video generation, 2025. 3
- 591 [29] Xiaofeng Mao, Shaoheng Lin, Zhen Li, Chuanhao Li, Wen-
592 shuo Peng, Tong He, Jiangmiao Pang, Mingmin Chi, Yu
593 Qiao, and Kaipeng Zhang. Yume: An interactive world gen-
594 eration model. *arXiv preprint arXiv:2507.17744*, 2025. 3
- 595 [30] Chenlin Meng, Yutong He, Yang Song, Jiaming Song, Jiajun
596 Wu, Jun-Yan Zhu, and Stefano Ermon. Sdedit: Guided image
597 synthesis and editing with stochastic differential equations.
598 *arXiv preprint arXiv:2108.01073*, 2021. 4
- 599 [31] Junho Park, Andrew Sangwoo Ye, and Taein Kwon.
600 Egoworld: Translating exocentric view to egocentric view
601 using rich exocentric observations. *arXiv preprint
602 arXiv:2506.17896*, 2025. 2
- 603 [32] Minho Park, Taewoong Kang, Jooyeon Yun, Sungwon
604 Hwang, and Jaegul Choo. Spherediff: Tuning-free omnidi-
605 rectional panoramic image and video generation via spheri-
606 cal latent representation. *arXiv preprint arXiv:2504.14396*,
607 2025. 3
- 608 [33] Nikhila Ravi, Jeremy Reizenstein, David Novotny, Daniel
609 Gordon, Wan-Yen Lo, Justin Johnson, and Georgia Gkioxari.
610 Pytorch3d: An open-source library for 3d deep learning. In
611 *CVPR Workshops*, 2020. 3
- 612 [34] Nikhila Ravi, Valentin Gabeur, Yuan-Ting Hu, Ronghang
613 Hu, Chaitanya Ryali, Tengyu Ma, Haitham Khedr, Roman
614 Rädle, Chloe Rolland, Laura Gustafson, Eric Mintun, Junting
615 Pan, Kalyan Vasudev Alwala, Nicolas Carion, Chao-
616 Yuan Wu, Ross Girshick, Piotr Dollár, and Christoph Feicht-
617 enhofer. Sam 2: Segment anything in images and videos.
618 *arXiv preprint arXiv:2408.00714*, 2024. 6
- 619 [35] Xuanchi Ren, Tianchang Shen, Jiahui Huang, Huan Ling,
620 Yifan Lu, Merlin Nimier-David, Thomas Müller, Alexan-
621 der Keller, Sanja Fidler, and Jun Gao. Gen3c: 3d-informed
622 world-consistent video generation with precise camera con-
623 trol. In *Proceedings of the Computer Vision and Pattern
624 Recognition Conference*, pages 6121–6132, 2025. 2, 3
- 625 [36] Oriane Siméoni, Huy V. Vo, Maximilian Seitzer, Federico
626 Baldassarre, Maxime Oquab, Cijo Jose, Vasil Khalidov,
627 Marc Szafraniec, Seungeun Yi, Michaël Ramamonjisoa,
628 Francisco Massa, Daniel Haziza, Luca Wehrstedt, Jianyuan
629 Wang, Timothée Darcret, Théo Moutakanni, Leonel Sentana,
630 Claire Roberts, Andrea Vedaldi, Jamie Tolan, John Brandt,
631 Camille Couprise, Julien Mairal, Hervé Jégou, Patrick La-
632 batut, and Piotr Bojanowski. DINOV3, 2025. 6
- 633 [37] Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Wen
634 Bo, and Yunfeng Liu. Roformer: Enhanced transformer with
635 rotary position embedding. *Neurocomputing*, 568:127063,
636 2024. 5
- [38] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30, 2017. 4
- [39] Team Wan, Ang Wang, Baole Ai, Bin Wen, Chaojie Mao, Chen-Wei Xie, Di Chen, Feiwu Yu, Haiming Zhao, Jianxiao Yang, et al. Wan: Open and advanced large-scale video generative models. *arXiv preprint arXiv:2503.20314*, 2025. 3, 5
- [40] Ruicheng Wang, Sicheng Xu, Yue Dong, Yu Deng, Jianfeng Xiang, Zelong Lv, Guangzhong Sun, Xin Tong, and Jiaolong Yang. Moge-2: Accurate monocular geometry with metric scale and sharp details. *arXiv preprint arXiv:2507.02546*, 2025. 3
- [41] Zun Wang, Jaemin Cho, Jialu Li, Han Lin, Jaehong Yoon, Yue Zhang, and Mohit Bansal. EPiC: Efficient Video Camera Control Learning with Precise Anchor-Video. *arXiv preprint arXiv:2505.21876*, 2025. 3
- [42] Zeqi Xiao, Yushi Lan, Yifan Zhou, Wenqi Ouyang, Shuai Yang, Yanhong Zeng, and Xingang Pan. Worldmem: Long-term consistent world simulation with memory. *arXiv preprint arXiv:2504.12369*, 2025. 3
- [43] Jilan Xu, Yifei Huang, Baoqi Pei, Junlin Hou, Qingqiu Li, Guo Chen, Yuejie Zhang, Rui Feng, and Weidi Xie. Egoexogen: Ego-centric video prediction by watching exo-centric videos. *arXiv preprint arXiv:2504.11732*, 2025. 2, 3, 5
- [44] Bowen Xue, Qixin Yan, Wenjing Wang, Hao Liu, and Chen Li. Stand-in: A lightweight and plug-and-play identity control for video generation. *arXiv preprint arXiv:2508.07901*, 2025. 3
- [45] Zhuoyi Yang, Jiayan Teng, Wendi Zheng, Ming Ding, Shiyu Huang, Jiazhen Xu, Yuanming Yang, Wenyi Hong, Xiaohan Zhang, Guanyu Feng, et al. Cogvideox: Text-to-video diffusion models with an expert transformer. *arXiv preprint arXiv:2408.06072*, 2024. 3
- [46] Deheng Ye, Fangyun Zhou, Jiacheng Lv, Jianqi Ma, Jun Zhang, Junyan Lv, Junyou Li, Minwen Deng, Mingyu Yang, Qiang Fu, et al. Yan: Foundational interactive video generation. *arXiv preprint arXiv:2508.08601*, 2025. 3
- [47] Mark YU, Wenbo Hu, Jinbo Xing, and Ying Shan. Trajectorycrafter: Redirecting camera trajectory for monocular videos via diffusion models. *arXiv preprint arXiv:2503.05638*, 2025. 2, 3, 5
- [48] Fan Zhang, Shulin Tian, Ziqi Huang, Yu Qiao, and Ziwei Liu. Evaluation agent: Efficient and promptable evaluation framework for visual generative models. *arXiv preprint arXiv:2412.09645*, 2024. 6
- [49] Lvmin Zhang, Anyi Rao, and Maneesh Agrawala. Adding conditional control to text-to-image diffusion models, 2023. 3